

Satellite Image Analysis with CNNs for Environmental Monitoring

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os
import gc

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Sequential
import keras as k
from keras import backend as K
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D, GlobalMaxPooling2D, BatchNormalization
import cv2
from tqdm import tqdm
from collections import Counter
from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split
from sklearn.metrics import fbeta_score
import plotly.express as px
```

Loading The Dataset

```
In [2]: path = "../input/planets-dataset/planet/planet/"
path_train = os.path.join(path, "train-jpg")
path_test = os.path.join(path, "test-jpg")
print( f"train files: {len(os.listdir(path_train))}, "
      f"test files: {len(os.listdir(path_test))}")
```

```
train files: 40479, test files: 40669
```

In [3]:

```
path_train_class = os.path.join(path, "train_classes.csv")
```

```
df_train = pd.read_csv(path_train_class)
```

```
print(df_train.shape)
```

```
df_train.head()
```

```
(40479, 2)
```

Out[3]:

	image_name	tags
0	train_0	haze primary
1	train_1	agriculture clear primary water
2	train_2	clear primary
3	train_3	clear primary
4	train_4	agriculture clear habitation primary road

Exploring and Understanding the Labels in the dataset

```
In [4]: all_tags = [item for sublist in list(df_train['tags'].apply(lambda row:
row.split(" ")).values) for item in sublist]
```

```
print('total of {} non-unique tags in all training images'.format(len(all_tags)))
```

```
print('average number of labels per image
```

```
{}'.format(1.0*len(all_tags)/df_train.shape[0]))
```

```
total of 116278 non-unique tags in all training images
```

```
average number of labels per image 2.8725511993873365
```

```
df_train["list_tags"] = df_train.tags.str.split(" ")
```

```
row_tags = df_train.list_tags.values
```

```
tags = [tag for row in row_tags for tag in row]
```

```
df_tags = pd.DataFrame(
```

```
    {"tag": counter_tags.keys(), "total": counter_tags.values()})
```

```
).sort_values("total")
```

```
fig = px.bar(df_tags, x="total", y="tag", orientation="h",color="total",)
```

```
fig.update_layout(title="Tags distribution")
```

```
fig.show()
```

Machine Learning

Preparing the Data

```
df_train = df_train.drop("list_tags", axis='columns')
df_train.head()
```

Out[6]:

	image_name	tags
0	train_0	haze primary
1	train_1	agriculture clear primary water
2	train_2	clear primary
3	train_3	clear primary
4	train_4	agriculture clear habitation primary road

```
x_train = []
y_train = []
flatten = lambda l: [item for sublist in l for item in sublist]
labels = list(set(flatten([l.split(' ') for l in df_train['tags'].values])))
label_map = {l: i for i, l in enumerate(labels)}
inv_label_map = {i: l for l, i in label_map.items()}
for f, tags in tqdm(df_train.values, miniters=1000):
    img = cv2.imread('../input/planets-dataset/planet/planet/train-
jpg/{f}.jpg'.format(f))
    targets = np.zeros(17)
    for t in tags.split(' '):
        targets[label_map[t]] = 1
    x_train.append(cv2.resize(img, (64, 64))) # Indicate the IMG Size
    y_train.append(targets)
x_train = np.array(x_train, np.float16) / 255.
y_train = np.array(y_train, np.uint8)
```

```
100%|██████████| 40479/40479 [07:41<00:00, 87.67it/s]
```

In [8]:

```
y_train = np.array(y_train, np.uint8)
x_train = np.array(x_train, np.float16) / 255.0
```

Splitting the data into train and validation sets.

```
x_train, x_val, y_train, y_val = train_test_split(x_train, y_train, test_size =
0.2, shuffle = True, random_state = 1)
```

```
# Prints the shape of the training and validation data.
```

```
print("Train data shape:",x_train.shape)
```

```
print("Train label shape:",y_train.shape)
```

```
print("Validation data shape:",x_val.shape)
```

```
print("Validation label shape:",y_val.shape)
```

```
Train data shape: (32383, 64, 64, 3)
```

```
Train label shape: (32383, 17)
```

```
Validation data shape: (8096, 64, 64, 3)
```

```
Validation label shape: (8096, 17)
```

```
gc.collect()
```

```
Out[9]:
```

```
96
```

Establishing Evaluation Metrics for the Model

In [10]:

```
def fbeta(y_true, y_pred, threshold_shift=0):  
    beta = 2  
  
    # Clipping y_pred between 0 and 1  
    y_pred = K.clip(y_pred, 0, 1)  
  
    # Rounding y_pred to binary values  
    y_pred_bin = K.round(y_pred + threshold_shift)  
  
    # Counting true positives, false positives, and false negatives  
    tp = K.sum(K.round(y_true * y_pred_bin)) + K.epsilon()  
    fp = K.sum(K.round(K.clip(y_pred_bin - y_true, 0, 1)))  
    fn = K.sum(K.round(K.clip(y_true - y_pred, 0, 1)))  
  
    # Calculating precision and recall  
    precision = tp / (tp + fp)  
    recall = tp / (tp + fn)  
  
def accuracy_score(y_true, y_pred, epsilon = 1e-4):  
    y_true = tf.cast(y_true, tf.float32)  
    y_pred = tf.cast(tf.greater(tf.cast(y_pred, tf.float32), tf.constant(0.5)), tf.float32)  
    tp = tf.reduce_sum(y_true * y_pred, axis = 1)  
    fp = tf.reduce_sum(y_pred, axis = 1) - tp  
    fn = tf.reduce_sum(y_true, axis = 1) - tp  
    y_true = tf.cast(y_true, tf.bool)  
    y_pred = tf.cast(y_pred, tf.bool)  
    tn = tf.reduce_sum(tf.cast(tf.logical_not(y_true), tf.float32) * tf.cast(tf.logical_not(y_pred), tf.float32),  
axis = 1)  
    return (tp + tn)/(tp + tn + fp + fn + epsilon)
```

Constructing the Neural Network Architecture

```
optimizer_Adam = Adam()
optimizer_Adagrad = Adagrad()
optimizer_RMSprop = RMSprop()
model = keras.Sequential()
model.add(BatchNormalization(input_shape=(64, 64, 3)))
model.add(Conv2D(32, kernel_size=(3, 3), padding='same', activation='relu'))
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.2))
model.add(Conv2D(64, kernel_size=(3, 3), padding='same', activation='relu'))
model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.2))
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(17, activation='sigmoid'))
model.compile(optimizer=optimizer_Adam,
              loss='binary_crossentropy',
              metrics=[fbeta, accuracy_score],
              history = model.fit(x_train, y_train,
                                batch_size=128,
                                epochs=10,
                                verbose=1,
                                validation_data=(x_val, y_val))
```

```
2023-01-19 05:34:51.633669: I tensorflow/core/common_runtime/process_util.cc:146] Creating new thread pool with default inter op setting: 2. To
use using inter_op_parallelism_threads for best performance.
```

2023-01-19 05:34:55.028825: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimization Passes are enabled (registered 2)

Epoch 1/10

253/253 [=====] - 231s 907ms/step - loss: 0.2369 - fbeta: 0.6281 - accuracy_score: 0.9103 - val_loss: 0.2538 - val_fb
eta: 0.6118 - val_accuracy_score: 0.9069

Epoch 2/10

253/253 [=====] - 227s 896ms/step - loss: 0.1970 - fbeta: 0.6946 - accuracy_score: 0.9231 - val_loss: 0.2217 - val_fb
eta: 0.6582 - val_accuracy_score: 0.9170

Epoch 3/10

253/253 [=====] - 230s 910ms/step - loss: 0.1728 - fbeta: 0.7379 - accuracy_score: 0.9323 - val_loss: 0.1644 - val_fb
eta: 0.7372 - val_accuracy_score: 0.9352

Epoch 4/10

253/253 [=====] - 230s 908ms/step - loss: 0.1617 - fbeta: 0.7569 - accuracy_score: 0.9370 - val_loss: 0.1504 - val_fb
eta: 0.7691 - val_accuracy_score: 0.9412

Epoch 5/10

253/253 [=====] - 231s 914ms/step - loss: 0.1526 - fbeta: 0.7725 - accuracy_score: 0.9408 - val_loss: 0.1439 - val_fb
eta: 0.7831 - val_accuracy_score: 0.9431

Epoch 6/10

253/253 [=====] - 229s 907ms/step - loss: 0.1456 - fbeta: 0.7820 - accuracy_score: 0.9429 - val_loss: 0.1397 - val_fb
eta: 0.7882 - val_accuracy_score: 0.9448

Epoch 7/10

253/253 [=====] - 227s 898ms/step - loss: 0.1413 - fbeta: 0.7901 - accuracy_score: 0.9447 - val_loss: 0.1371 - val_fb
eta: 0.7848 - val_accuracy_score: 0.9454

Epoch 8/10

253/253 [=====] - 228s 901ms/step - loss: 0.1372 - fbeta: 0.7963 - accuracy_score: 0.9462 - val_loss: 0.1354 - val_fb
eta: 0.7900 - val_accuracy_score: 0.9457

Epoch 9/10

253/253 [=====] - 231s 911ms/step - loss: 0.1321 - fbeta: 0.8030 - accuracy_score: 0.9476 - val_loss: 0.1337 - val_fb
eta: 0.8061 - val_accuracy_score: 0.9470


Epoch 10/10

253/253 [=====] - 228s 901ms/step - loss: 0.1275 - fbeta: 0.8117 - accuracy_score: 0.9493 - val_loss: 0.1335 - val_fb
eta: 0.7949 - val_accuracy_score: 0.9474

plt.plot(history.history['loss'])

plt.plot(history.history['val_loss'])

```
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



```
model.evaluate(x_val, y_val)
train_fscore = fbeta_score(y_train, np.round(model.predict(x_train)), beta=2, average = 'weighted')
print("train fscore: ", train_fscore)
val_fscore = fbeta_score(y_val, np.round(model.predict(x_val)), beta=2, average = 'weighted')
print("val fscore: ", val_fscore)
```

```
253/253 [=====] - 15s 59ms/step - loss: 0.1335
- fbeta: 0.7952 - accuracy_score: 0.9474
train fscore: 0.7965989848478681
val fscore: 0.7771786500700735
```



```

df_samplesub = pd.read_csv('../input/planets-
dataset/planet/planet/sample_submission.csv')
test = df_samplesub[0 : 40669]

files = df_samplesub[40669 : ]

test_img = []

for image_name, tags in tqdm(test.values, miniters=1000):
arr = cv2.imread('../input/planets-dataset/planet/planet/test-jpg/{0}.jpg'.format(image_name))
test_img.append(cv2.resize(arr, (64, 64)))

for image_name, tags in tqdm(files.values, miniters=1000):
arr = cv2.imread('../input/planets-dataset/test-jpg-additional/test-jpg-
additional/{0}.jpg'.format(image_name))
test_img.append(cv2.resize(arr, (64, 64)))

test_img = np.array(test_img, np.float16)/255.0

```

```

100%|██████████| 40669/40669 [08:17<00:00, 81.73it/s]
100%|██████████| 20522/20522 [03:52<00:00, 88.44it/s]

```

```
gc.collect()
```

```

Out[19]:
6626

```

```

yres = []
predictions = model.predict(test_img, batch_size = 64, verbose = 2)
yres.append(predictions)
957/957 - 102s
gc.collect()

```

```

Out[21]:
788

```

```

sub = np.array(yres[0])
for i in range(1, len(yres)):

```

```

sub += np.array(yres[i])
sub = pd.DataFrame(sub, columns = label_map)
preds = []

# Loop through the sample submission DataFrame
for i in tqdm(range(sub.shape[0]), miniters=1000):

    # Get the i-th row of the DataFrame
    a = sub.loc[[i]]

    # Apply a lambda function to get a Boolean array indicating which columns have values greater than 0.2
    a = a.apply(lambda x: x > 0.2, axis=1)

    # Transpose the DataFrame
    a = a.transpose()

    # Get the rows where the Boolean array is True
    a = a.loc[a[i] == True]

    # Join the index of the DataFrame (which contains the tags) into a single string
    tag = ''.join(list(a.index))

    # Append the string of tags to the preds list
    preds.append(' '.join(list(a.index)))

# Assign the preds list as the 'tags' column of the sample submission DataFrame
df_samplesub['tags'] = preds

# Save the sample submission DataFrame to a CSV file
df_samplesub.to_csv('CMT_submission.csv', index=False)
100%|██████████| 61191/61191 [01:51<00:00, 550.94it/s]

```

image_nar tags

test_0 cloudy

test_1 cloudy

test_2

test_3 cloudy

test_4

test_5 cloudy

test_6 cloudy

test_7 primary water clear artisinal_mine

test_8 cloudy

test_9 cloudy